

CSE13S notes

Week of 1/30/23

Assignment 5: Schmidt-Samoa cryptography

- GMP has a function to generate random numbers
- Write utility functions - small ones can still be efficient since compiler will just paste the small function where you've called it
- Miller-prime works better the more iterations of it you go through
- All about the algorithm - code in pure C first before using GMP

Compilers!! (Innards of language translators)

- Simply put, a program that maps an input language to an output language (much easier to do programming language -> assembly but much more difficult to translate human language)
- Roles:
 - Convert high-level source code to machine language
- In the good old days:
 - Computers were programming in binary
 - Hexidecimal
 - Octal
 - It was hard
 - If you hated assembly in cse12, well this was much harder

Why need language translators?

- Computers are a balanced mix of software and hardware
- Hardware functions are controlled by a compatible software program

Compilation process of a C program

- The compiling process is a sequence of four phases. The input for each phase is the output of the previous
 - Source code -> preprocessor -> compiler -> assembler -> linker - exe
- The pre-processor
 - First thing that runs, usually removes comments
 - Pastes in included files and standard library headers
- The compiler
 - The compiler converts hello.i from the preprocessor into assembly code in 7 steps:
 - Lexical
 - Syntax
 - Translation

- Optimization
 - Storage assignment
 - Code generation
 - Assembly phase
 - Lexical phase breaks down source code into lexemes (see slides)
- The Assembler
 - Assembly code to machine code (binary) is very simple 1 to 1 translation
 - Can perform conversions is 2 ways - (see slides)
 - Forward referencing happens in one or two passes (see slides)
 -
- The linkers
 - Links hello.o along with any other object files and libraries
 - Ensures all dependencies are properly resolved
 - Linker error thrown if definitions to something that cannot be referenced
 - Merges everything into one executable file
 - In this case, hello.exe
- Loaders
 - Lays in the operating system
 - Ensures program and necessary libraries are placed in RAM to prepare them for execution
 - The loader has 4 basic functions
 - Allocation: allocate space in memory for program
 - linking - resolve symbolic references between programs
 - Relocation - fix all dependent locations and point them to the newly allocated space
 - Loading - place machine code and data directly into the processor
- Memory layout
 - See slides for memory layout in unix
- Compilers vs interpreters
 - Compiler
 - A translator that takes a high-level programming language, goes through a sequence of translations and output an executable
 - Translates entire program/code at once
 - Interpreter
 - Directly executes code without needing to compile (python interpreter, ghci (haskell))
 - Translates program one line at a time
 - Which is faster, the compiler or an interpreter? Generally, the compiler

Interoperability

- Code should be interoperable between different environments (as much as is plausible)

Files

Long-term information storage

- We want to store large amounts of data
 - The definition of large has changed by 6 orders of magnitude over the years
- Information stored must survive the termination of the process using it
 - Memory (DRAM) does not survive turning the computer off
 - This is why you don't kill the VM to close it, it corrupts data because all memory is gone
- Files are accessed using names, memory is accessed using addresses (pointers/variables)

File names

- Our computers have millions of files
 - Data center has billions of files
 - The internet has trillions of files
- How do you uniquely specify which file you want?
- How do you find that file among millions of files?
- Files have base names and extensions (base.extension)
 - In some OSs, the extension is a separate entity
- In UNIX, the extension is only used by convention

Files may have internal structure

- .EXE have headers containing information about the file, as well as text (code), data, relocation bits, and symbol table
- Archive files have headers and object modules

File Access

- Sequential access
 - Read all bytes/records from the beginning
 - Cannot jump around, may rewind or back up
 - Convenient when medium was magnetic tape
- Random access
 - bytes/records read in any order
 - Essential for data base systems
 - Read can be:

- Move file marker (seek), then read or...
 - Read and then move file marker
- File attributes
 - Files have many attributes that can be put in a file (in windows and mac)
 - UNIX also has same attributes, but does not care about certain attributes
- File operations
 - Create, delete, open, close, read, write
 - Append, seek, get attributes, set attributes
 - Rename

Directories

- Naming is better than using numbers, but still limited
- Humans like to group things together for convenience
 - We use hierarchy to manage complexity
- File system allow this to be done with directories
 - You may call them folders
- Grouping makes it easier to:
 - Find files in the first place: remember the enclosing directories for the file,
 - Locate related files,
 - Determine which files are related
- Directory operations
 - Create, delete, opendir, closedir
 - Readdir, rename, link, unlink

2/17/23 notes

Processes

- What is a process? *A running program*
- code , data, and stack have its own address space
- Program state (what the state of the program is right now)
 - CPU registers
 - Program counter (current location in the code)
 - Stack pointer
- Only one process can be running in a single CPU core at any given time
 - Multi-core CPUs can support multiple processes

What is an Address space?

- Program execute code
 - Each instruction has an address
- Programs access data

- Each byte of data also has an address
- We would like to think that our program is the only program executing on the computer
 - But we would be wrong
- See slides

How is that accomplished (what is that?)

- The loader could relocate the instructions by address a base address to each one
- There could be registers that point to the first byte of the program's memory (base register) that is added to every memory access
- The processor and operating system could provide virtual memory

In an ideal world...

- The ideal world has memory that is
 - Very large, very fast, and non-volatile
 - But we don't live in an ideal world
- The real world has memory that is (PICK 2):
 - Large
 - Fast
 - Affordable
- The goal of memory management is to reach for an ideal world as much as possible

Memory hierarchy

- What is the memory hierarchy?
 - Different levels of memory
 - Some are small & fast
 - Others are large & slow
- What levels are usually included?
 - Cache: small amount of fast, expensive memory
 - L1 (level 1) cache: usually on the CPU chip
 - L2: may be on or off chip
 - L3 cache: off-chip, made of SRAM
 - Main memory: medium-speed, medium price memory (DRAM)
 - Disk: many gigs of slow, cheap, non-volatile storage
- Memory manager handles the memory hierarchy

Basic memory management

- Components include
 - Operating system (perhaps with device drivers)
- Fixed partitions: multiple programs
 - Fixed memory partitions created, where memory is divide into fixed spaces
 - Assign a process to a space when it's free
- Mechanisms
 - Separate input queues for each partition
 - Single input queue: better ability to optimize CPU usage
- Multiprocessing works because I/O takes a long long time
 - Most programs really are waiting for I/O, so cpu can work on other programs

Memory and multiprogramming

- See slides

Base and limit registers

- See slides

Allocating memory

- Search through region list to find a large enough space
- Suppose there are several choices: which one to use?
 - First fit: the first suitable hole on the list
 - Next fit: the first suitable after the previously allocated hole
 - Best fit: the smallest hole that is larger than the desired region (wastes least space?)
 - Best is worst because of the small, unusable memory fragments that get left in the memory space
 - Worst fit: the largest hole that is available

Freeing memory

- Allocated structures must be updated when memory is freed
- Easy with bitmaps: just set the appropriate bits in the bitmap
- Linked lists: modify adjacent elements as needed
 - Merge adjacent free regions into a single region
 - May involve merging two regions with the just-freed area

Buddy allocation

- Allocates memory into powers of 2
 - Partitions in half until a suitably small region is found for an object
- Split larger chunks to create 2 smaller regions

Virtual memory

- Basic idea: allow the OS to hand out more memory than exists on the system

- Keep recently used stuff in physical memory
- Move less recently used stuff to disk
- Keep all of this hidden from processes
 - Processes still see an address space from 0-max_address
 - Movement of information to and from disk handled by the OS without process help
- Virtual memory (VM) especially helpful in multiprogrammed systems
 - CPU schedules process B while process A waits for its memory to be retrieved from the disk

Virtual and physical addresses

- Program uses virtual addresses
 - Addresses local to the process
 - Hardware translates virtual address to physical address
- Translation done by the memory management unit
 - Usually on the same chip as the CPU
 - Only physical addresses leave the CPU/MMU chip
- Physical memory indexed by physical addresses
- Virtual memory can extend anywhere on physical memory, so it is better to use (prevents small unusable memory spaces)

What's in a page table entry?

- Each entry in the table entry?
 - Protection: who can read/modify that piece of memory
 - Dirty (modified) bit: set if data in the page has been modified
 - Reference bit:
 - Valid bit:
- See slides

The process model

- Conceptual model of multiprogramming four programs
 - 4 independent processes
 - Processes run sequentially
- Only one process active at any time

When is a process created?

- Processes can be created in two ways:
 - System initialization: one or more processes created when the OS starts up
 - Execution of a process creation system call: something explicitly asks for a new process
- System call comes from (see slides)

When do processes end?

- Conditions that terminate processes that do

Processes also have a process table counter

-